


M=8: layer-2 attention (one example, $a = 4$)

legible: a sharp spike on m_4



M=32: layer-3 attention (averaged over test set)

illegible: diffuse across memory

